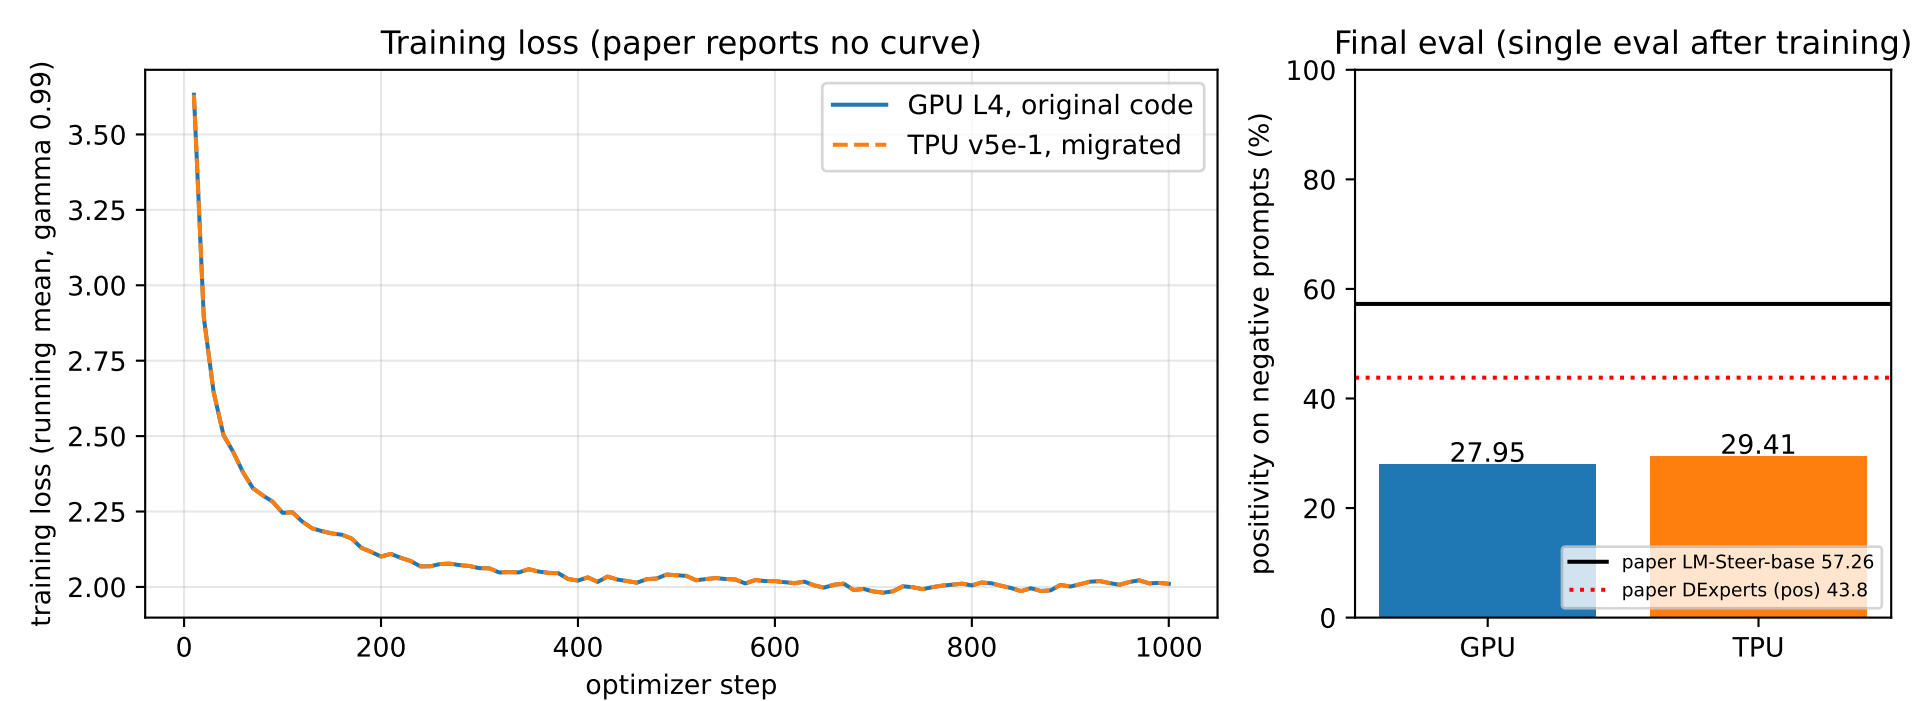
Training loss (paper reports no curve)
training loss (running mean, gamma 0.99)
optimizer step
GPU L4, original code
TPU v5e-1, migrated
Final eval (single eval after training)
positivity on negative prompts (%)
27.95
29.41
GPU
TPU
paper LM-Steer-base 57.26
paper DExperts (pos) 43.8